

Hunting for Pulsars

Joe Davies

Useful Links

- Particle physics application:
<https://ilmonteux.github.io/2018/10/15/jet-tagging-cnn.html>
- Kaggle page for the pulsar dataset:
<https://www.kaggle.com/shivam1901/pulsar-star>
- Github repo with all the code from these tutorials:
<https://github.com/adrianbevan/IntroToML/tree/master/Pulsars>
- <https://learn.datacamp.com/courses/introduction-to-deep-learning-with-keras>
- Misc useful links for more information on machine learning:
 - <https://towardsdatascience.com/>
 - <https://www.coursera.org/learn/machine-learning>
 - <https://elitedatascience.com/learn-machine-learning>
- j.m.m.davies@qmul.ac.uk

Introductions!




The Dark Machines Anomaly Score Challenge: Benchmark Data and Model Independent Event Classification for the Large Hadron Collider

T. Aarrestad^{CERN} M. van Beekveld^{Ox} M. Bona^{QMUL} A. Boveia^{OSU}
S. Caron^{HEF, Nikhef} J. Davies^{QMUL} A. De Simone^{SISSA, INFN} C. Doglioni^{Lund}
J.M. Duarte^{UCSD} A. Farbin^{UnivArlington} H. Gupta^{GSoC} L. Hendriks^{HEF, Nikhef}
L. Heinrich^{CERN} J. Howarth^{Glasgow} P. Jawahar^{WPI, CERN} A. Jueid^{UnivKonkuk}
J. Lastow^{Lund} A. Leinweber^{UnivAdelaide} J. Mamuzic^{IFIC} E. Merényi^{UnivRice}
A. Morandini^{RWTH} P. Moskvitina^{HEF, Nikhef} C. Nellist^{HEF, Nikhef}
J. Ngadiuba^{FNAL, Caltech} B. Ostdiek^{Harvard, AIFI} M. Pierini^{CERN} B. Ravina^{Glasgow}
R. Ruiz de Austri^{IFIC} S. Sekmen^{KNU} M. Touranakou^{NKUA, CERN}
M. Vaškevičiūtė^{Glasgow} R. Vilalta^{UnivHouston} J.-R. Vlimant^{Caltech} R. Verheyen^{UCL}
M. White^{UnivAdelaide} E. Wulff^{Lund} E. Wallin^{Lund} K.A. Wozniak^{UniVie, CERN}
Z. Zhang^{HEF, Nikhef}




Machine Learning vs Artificial Intelligence





<https://www.edureka.co/blog/ai-vs-machine-learning-vs-deep-learning/>

The Data

- Data is based on a Kaggle set investigating pulsars
- Pulsars are highly magnetized neutron stars that emit radiation from their magnetic poles
- Data contains 8 columns including metrics like kurtosis and dispersion of radiation measures
- The data also contains a target class: 0 or 1 depending on not-a-pulsar or pulsar



pulsar_stars.csv (1.67 MB)															9 of 9 columns		View	
#	Mean of the integrated profile.	#	Standard deviation of the integrated profile.	#	Excess kurtosis of the integrated profile.	#	Skewness of the integrated profile.	#	Mean of the DM-SNR curve.	#	Standard deviation of the DM-SNR curve.	#	Excess kurtosis of the DM-SNR curve.	#	Skewness of the DM-SNR curve.	#	target_class	
1. Mean of the integrated profile.																		9 classes of pulsar star. 1 for pulsar star,0 for not a star
																		
5.81 193 24.8 98.8 1.88 8.07 -1.79 68.1 0.21 223 7.37 111 -3.14 34.5 -1.98 119k 0																		
1	140.5625	55.68378214	-0.234571412	-0.699648398	3.199832776	19.11842633	7.975531794	74.24222492	€									
2	102.5078125	58.88243001	0.465318154	-0.515887909	1.677257525	14.86014572	10.57648674	127.3935796	€									
3	103.815625	39.34164944	0.323328365	1.851164429	3.121237458	21.74466875	7.735822015	63.17190911	€									
4	136.75	57.17844874	-0.868414638	-0.636238369	3.642976589	20.9592883	6.89649891	53.59360667	€									
5	88.7265625	40.67222541	0.600666079	1.123491692	1.178929766	11.4687196	14.26957284	252.5673858	€									
6	93.5789125	46.69811352	0.53190485	0.416721117	1.636287625	14.54587425	16.6217484	131.3940043	€									
7	119.484375	48.76585927	0.831468022	-0.112167573	0.99916388	9.279612239	19.20623018	479.7565669	€									
8	130.3828125	39.84485561	-0.158322759	0.389548448	1.220735786	14.37894124	13.53945602	198.2364565	€									
9	107.25	52.62707834	0.452688025	0.178347382	2.331939799	14.48685311	9.801004441	107.9725856	€									
10	107.2578125	39.49648839	0.465881961	1.162877124	4.079431438	24.98041798	7.397879948	57.78473789	€									
11	142.878125	45.28987262	-0.320328426	0.283952506	5.376254101	29.80989748	6.876265849	37.83139335	€									
12	133.2578125	44.05824378	-0.881059862	0.115361506	1.632107023	12.80788568	11.97206663	195.5434476	€									
13	134.9689375	49.55432662	-0.135803833	-0.880496062	10.69648829	41.34284361	3.893934139	14.13128625	€									
14	117.9453125	45.50657724	0.325437564	0.661459458	2.836120401	23.11834971	8.943211912	82.47559187	€									
15	138.1796875	51.5244835	-0.831852329	0.846797173	6.380267559	31.57634673	5.155939859	26.14331017	€									
16	114.3671875	51.94571552	-0.894498904	-0.287984087	2.738294314	17.19189079	9.850612454	96.61190318	€									
17	109.648625	49.01765217	0.13763583	-0.25669775	1.508361204	12.87290134	13.36792556	223.4384192	€									
18	100.8515625	51.74352161	0.393836792	-0.811240741	2.841137124	21.63577754	8.302241891	71.58436903	€									
19	136.89375	51.69100464	-0.845908926	-0.271816393	9.342809365	38.89639955	4.345438138	18.67364854	€									
20	99.3671875	51.57220208	1.547196967	4.154106843	27.55518395	61.71901588	2.20880796	3.662800136	1									
21	100.890625	51.89039446	-0.826497802	-0.627486528	3.883779264	23.84526673	6.953167635	52.27944038	€									
22	105.4453125	41.13966851	0.142653801	0.320419676	3.551839465	20.75581684	7.739552295	68.51977061	€									
23	95.8671875	42.05992212	0.326386917	0.803581794	1.83277592	12.24896949	11.249331	177.2307712	€									
24	117.3671875	53.90861351	0.257953441	-0.485849077	6.018394649	24.76612335	4.807783224	25.52261561	€									

What are we doing?

- **Create 3 machine learning algorithms that identify pulsars from the data**
- **Understand how to use specific python modules to do this**
- **Get an idea of when to use and not use each algorithm**
- **sklearn, scipy, pandas, matplotlib, numpy, keras, tensorflow**


Logistic Regression

- Statistical model that uses a logistic function to (usually) model a binary variable (0 or 1)

- An example of this function is the sigmoid:

$$f(x) = \frac{1}{1 + e^{-x}}$$

- Uses gradient descent for the meat of the regression
- An issue with this approach is that extreme cases are presumed to become progressively rarer at a specific rate
 - Not very good if one doesn't have much data



Logistic Regression

Example file: pulsar_logistic_sklearn.py

Data Preparation

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from scikitplot.estimators import plot_learning_curve
df=pd.read_csv('~/.Documents/ML/pulsars/pulsar_stars.csv')
```

```
x_data=df.drop(columns='target_class')
x=(x_data-np.min(x_data))/(np.max(x_data)-np.min(x_data)) #scaling
y=df.target_class.values

compare_score=[]

#training and testing split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=123)
```

Model Creation


```
lr=LogisticRegression()
lr.fit(x_train, y_train)
```

Visualization



```
lr_score=lr.score(x_test, y_test) * 100
compare_score.append(lr_score)

print('Test accuracy: {}'.format(lr_score))

plot_learning_curve(lr, x_test, y_test)
plt.show()
```



Decision Trees





Advantages

- Simple to understand and interpret
- Useful for classification and regression
- Requires little data preparation and can handle large data sets

Disadvantages

- Not very robust
- Make locally optimal decisions (info gain max, entropy min)
- Prone to over-fitting

Boosted Decision Tree



- Boosting involves running multiple trees one after the other and minimizing error
- Each tree output, $h(x)$, is given a weight, w , based on accuracy: $\hat{y}(x) = \sum_t w_t h_t(x)$
- Each data sample is given a weight based on misclassification
- We minimize the objective function: $O(x) = \sum_i l(\hat{y}_i, y_i) + \sum_t \Omega(f_t)$
- First term is loss function, second is regularization (penalizes complexity of the t -th tree)

<https://towardsdatascience.com/understanding-adaboost-2f94f22d5bfe>

<https://indico.fnal.gov/event/15356/contributions/31377/attachments/19671/24560/DecisionTrees.pdf>

https://scikit-learn.org/stable/auto_examples/ensemble/plot_adaboost_regression.html

Boosted Decision Tree

Example file: pulsar_bdt.py

Data Preperation

```
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.model_selection import train_test_split
from sklearn import metrics
#Vizualization
import matplotlib.pyplot as plt
from scikitplot.estimators import plot_learning_curve

data=pd.read_csv('~/.Documents/ML/pulsars/pulsar_stars.csv')
```

```
features=data.columns[:-1]
X=data[features]
#output:
y=data.target_class

#Now we need to split the data using train_test_split
X_train, X_test, y_train, y_test=train_test_split(X, y,
                                                    test_size=0.2, random_state=42)
```

Model Creation


```
classifier=AdaBoostClassifier(DecisionTreeClassifier())

classifier=classifier.fit(X_train, y_train)
#Predicting the response for the test dataset
y_pred=classifier.predict(X_test)
```

Visualization


```
score=round(metrics.accuracy_score(y_test, y_pred)*100, 2)
print('Accuracy = {}'.format(score))

plot_learning_curve(classifier, X_test, y_test)
plt.show()
```



Artificial Neural Networks

$$y = f\left(\sum_{i=1}^N w_i x_i + \theta\right)$$
$$= f(w^T x + \theta)$$



$$w_{r+1} = w_r - \gamma \sum_{i=1}^N \frac{\partial E_i}{\partial w}$$
$$\theta_{r+1} = \theta_r - \gamma \sum_{i=1}^N \frac{\partial E_i}{\partial \theta}$$

Advantages

Among the most accurate of modelling approaches

Useful for classification and regression

Makes few assumptions about relationships in the data

Disadvantages

Computationally intensive

Easy to over- or under-training data:

Results in a complex black box model

Artificial Neural Network

Example file: pulsar_ann.py

Data Preperation

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from keras import Sequential
from keras.layers import Dense
from sklearn.metrics import confusion_matrix
from scikitplot.estimators import plot_learning_curve
df=pd.read_csv('~/.Documents/ML/pulsars/pulsar_stars.csv')
```

```
x_data=df.drop(columns='target_class')
X=StandardScaler().fit_transform(x_data)
y=df.target_class.values

x_train, x_test, y_train, y_test=train_test_split(X, y,
                                                    test_size=0.2, random_state=42)
```

Model Creation

```
classifier=Sequential()

#first hidden layer
#we have 8 input features, 1 output and the kernel_initializer uses a normal distribution to
#function
classifier.add(Dense(8, activation='relu', kernel_initializer='random_normal', input_dim=8))

#second
classifier.add(Dense(8, activation='relu', kernel_initializer='random_normal'))

#output
classifier.add(Dense(1, activation='sigmoid', kernel_initializer='random_normal'))

#compiling the network
classifier.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

#fitting the data to the training set
history=classifier.fit(x_train, y_train, validation_split=0.33, batch_size=10, epochs=15)

#evaluate the loss value and metrics values for the model in test mode using evaluate funcn.
eval_model=classifier.evaluate(x_train, y_train)
print('eval model: ', eval_model)
```

- 1 input layer, 1 hidden layer, 1 output layer
- 8 features, 1 output
- Minimum 8 neurons in the hidden layer
- Validation set to get an idea of accuracy per epoch

Artificial Neural Network


Example file: pulsar_ann.py

Visualization


```
cm=confusion_matrix(y_test, y_pred)
print(cm)

# list all data in history
print(history.history.keys())
# summarize history for accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```




```
[[3233  26]
 [ 42 279]]
```



Something cool to end: GANs



Backup: Gradient Descent



- $F(x)$ is differentiable around point θ
- F will decrease fastest if going in the direction of negative gradient
- γ is some real, positive number
- Converges to a local minimum

$$\theta_{i+1} = \theta_i - \gamma \nabla F(\theta_i)$$